

Task I

Pre-processing:

For preprocessing, we essentially flatten the input, convert it to lower - case and store all the words, we construct a word level tokenizer storing all the words mapped to their respective embeddings and an inverse map for embedding to word mappings. The tokenizer also consists of special tokens like [START], [END], [PAD], [UNK]

Model Architecture:


```
TransformerLM(  
    (token_embedding): Embedding(11641, 128)  
    (positional_encoding): PositionalEncoding(  
        (pos_embedding): Embedding(512, 128)  
    )  
    (dropout): Dropout(p=0.2, inplace=False)  
    (decoder_layers): ModuleList(  
        (0): TransformerBlock(  
            (self_attn): MultiHeadAttention(  
                (W_q): Linear(in_features=128, out_features=128, bias=True)  
                (W_k): Linear(in_features=128, out_features=128, bias=True)  
                (W_v): Linear(in_features=128, out_features=128, bias=True)  
                (W_o): Linear(in_features=128, out_features=128, bias=True)  
            )  
            (feed_forward): PositionWiseFeedForward(  
                (linear1): Linear(in_features=128, out_features=256, bias=True)  
                (linear2): Linear(in_features=256, out_features=128, bias=True)  
            )  
            (norm1): LayerNorm((128,), eps=1e-05, elementwise_affine=True)  
            (norm2): LayerNorm((128,), eps=1e-05, elementwise_affine=True)  
            (dropout): Dropout(p=0.2, inplace=False)  
        )  
    )  
    (fc_out): Linear(in_features=128, out_features=11641, bias=True)  
)
```

We use the nn.Embedding layer to create token embeddings, then we use a positional encoding layer to take the position of the tokens into account. Then we use a dropout layer to perform regularisation. We use 1 Decoder Style Transformer Block consisting of a Multi Head Attention block with 2 heads, a Position Wise FFN, Layer Norms and a dropout layer. And in the end we use a fully connected layer to generate token probabilities.

Hyper-Parameters:

- 1.Vocab Size - 13393
- 2.Embedding Dimension - 128
- 3.Number of Heads - 2
- 4.Number of Layers - 1
- 5.Dimension of FFN - 256
- 6.Dropout - 0.2
- 7.Learning Rate - $2e-3$
- 8.Weight Decay - $1e-4$
- 9.Label Smoothing - 0.01
10. Number of Epochs - 10

Training and Validation Loss Plots:



As we can see, the model starts overfitting after the 2nd Epoch.

Perplexity Score on validation sets:

100%|██████████| 308/308 [00:31<00:00, 9.89it/s]

Epoch 2: Train Loss = 5.2844, Val Loss = 5.3884, Perplexity Score Train = 197.2354, Preplexity Score Validation = 218.8594
Sample text: thou art is the most beautiful of my poor good heart , that that she will not a man in his name ?
: i hear for heaven 's right ; and his majesty .
with

Task II

Pre-Processing

An acronym and contraction dictionary was made by going through the dataset beforehand

Steps:

- All the characters are converted to lowercase letters
- Contractions library is used to expand any common contractions and slangs used in language
- Regex is used to find words which belong to the expansion dictionary and are replaced by their expansions
- Regex is used to remove all strings which start with www. Http https , removing almost all links
- Regex is used to remove any special character except for numbers and alphabets
- Regex is used to remove extra whitespace

Pre processed dataset is then stored as a csv file which can be loaded in further training

Sample function from data frames is used to split the data into 70% train , 15% val , 15% test

Model Architecture

T5

T5 base is used for tokenization and for generation


Before fine tuning the dataset was preprocessed by adding the word normalize before the input text , essentially giving it input as Normalize <input text> , and the label as the given claim all padded to the same lengths

Training:

- For training the trainer class was used
- The learning rate was set to $3e-4$

- Model was trained for 10 epochs
- Weight decay of 0.1 was added to prevent overfitting
- Batch size of 8 was used
- The best model was loaded for evaluation at the end of training
- Data collater was used to pad input and labels to the same length

Train and Val Loss plots for T5



Evaluation Metrics for T5

```

ROUGE-L Precision: 0.4087
ROUGE-L Recall:    0.4325
ROUGE-L F1 Score:  0.3975
BLEU Score:        0.2125
BERT Score F1:     0.8901

```

Bart

Bart large is used for tokenization and for generation


Before fine tuning the dataset was preprocessed by tokenizing and padding

Training:

- For training the trainer class was used
- The learning rate was set to $3e-5$
- Model was trained for 10 epochs

- Weight decay of 0.1 was added to prevent overfitting
- Batch size of 8 was used
- The best model was loaded for evaluation at the end of training
- Data collater was used to pad input and labels to the same length

Train and Val Loss plots for Bart



Evaluation Metrics for Bart

```

ROUGE-L Precision: 0.4855
ROUGE-L Recall:    0.4419
ROUGE-L F1 Score:  0.4400
BLEU Score:        0.2727
BERT Score F1:     0.9011

```

Comparative analysis

ROUGE-L Precision & F1: BART-large shows noticeably higher precision (0.4855 vs. 0.4087) and F1 (0.4400 vs. 0.3975), indicating that its summaries capture a higher proportion of the important content relative to the generated tokens.

BLEU Score: A BLEU score improvement (0.2727 vs. 0.2125) for BART-large suggests that its generated text matches reference texts more closely at the n-gram level.

BERT Score F1: With a score of 0.9011 compared to T5-base's 0.8901, BART-large achieves better semantic similarity between its outputs and the reference normalized claims.

Overall Bart outperforms T5

Resource Constraints and Model Selection:

When selecting models for fine-tuning on the claim normalization task, a key consideration was the balance between performance and available computational resources.

- T5-base offers a lower parameter count and reduced memory requirements, making it an attractive option when GPU memory and training time are limited.
- As the Dataset was limited going for a bigger model would have probably overfit and given the GPU requirements would have taken a lot longer to train
- Bart large was a bigger model but the training time was not significantly more , with the evaluation metrics being better for Bart large it was chosen

Task III

- Importing necessary libraries

Python libraries:

- os, numpy, matplotlib, pandas,
- torch, pickle, torchvision, csv, PIL,
- transformers, tqdm, evaluate & nltk

are imported

- Loading train and validation:

- Data Frames,
- Image Descriptions
- Objects Detected

are loaded

- Instances of

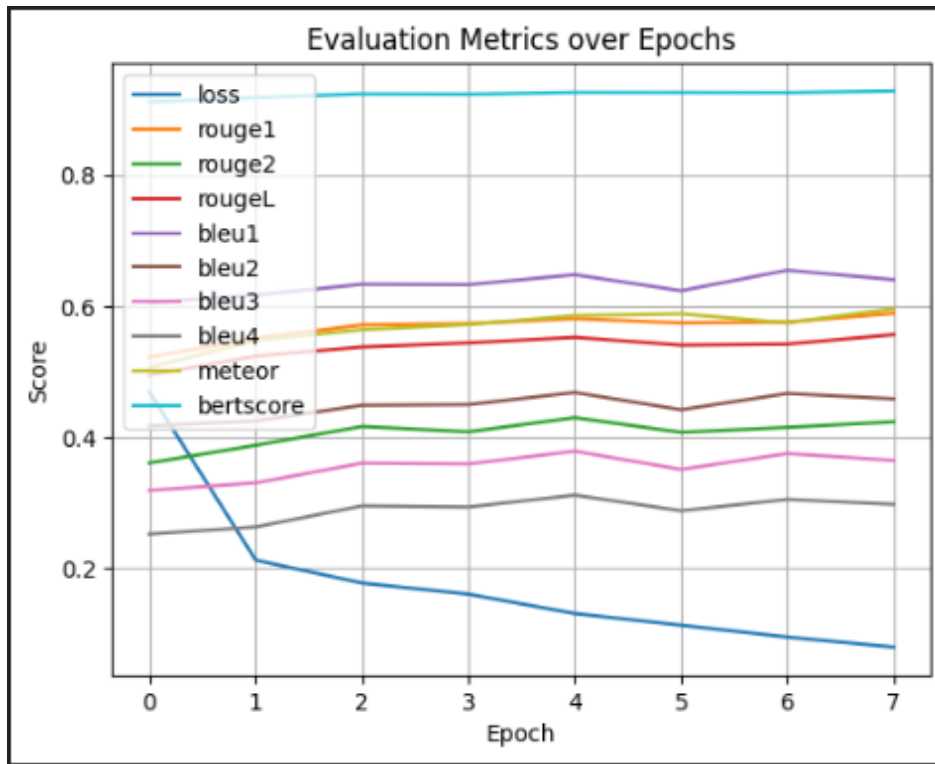
- Tokenizer
- BART Model
- ViT Model
- ViT Processor

are created from pre-trained methods

- Define the custom dataset & dataloader instances of train and validation
- Initialise *rouge*, *bleu*, *metoer* & *bertscore* instances for metric calculation. And initialise the metric-history dictionary
- Define TURBOModule, following the architecture mentioned in the paper. Notably, Knowledge Gates and GCNs are omitted as directed.
- Create a TURBOModule instance using BART and ViT Model instances.
- Create an instance of AdamW optimizer with *learning_rate* of **3e-5**
- Define method for plotting the metrics
- Define the training loop:
 - set checkpoint directory
 - training is looped over *num_epochs* times
 - at every epoch, each batch is iterated over, and model is used to extract predictions
 - Then, loss is calculated and a *backward* step is initialised
 - Model is saved for the epoch
 - The trained model is then evaluated over the *validation* dataset and metrics are calculated, with *metrics_history* is updated.

Additionally, based on the 'RougeL' score, the best performing model is updated. Finally, in a case where the model doesn't improve after *early_stopping_patience*, the model stops learning.

- Plot the metrics



Training loss: 0.0797

Validation Metrics:

- ROUGE-1: 0.5888
- ROUGE-2: 0.4235
- ROUGE-L: 0.5564
- BLEU-1 to BLEU-4: [0.6403, 0.4584, 0.3645, 0.2975]
- METEOR: 0.5961
- BERTScore (F1): 0.9282

- Hyperparameters:

- learning rate = $3e-5$
- maximum length = 128
- hidden dimension = 768

- Sample Validation predicting:

Example Predictions:

Original Post : '<user> thank u for this awesome network in malad (see pic) . # pathetics'
 [SEP] Sarcasm Target: <user>'s network in malad [SEP] Image Description: [SEP] Detected Objects:
 Ground Truth Exp : the author is pissed at <user> for not getting network in malad.
 Predicted Exp : the author is pissed at <user> for such awful network in malad.

Original Post : Nothing like waiting for an hour on the tarmac for a gate to come open in snowy, windy Chicago! [SEP] Sarcasm Target: gate not opening [SEP] Image Description: [SEP] Detected Objects:
Ground Truth Exp : nothing worst than waiting for an hour on the tarmac for a gate to come open in snowy, windy Chicago.
Predicted Exp : nothing like waiting for an hour on the tarmac for a gate to come not opening in snowy, windy Chicago.

Original Post : 'ahhh ! my favorite thing about spring ! # dst # springforward ' [SEP] Sarcasm Target: spring weather [SEP] Image Description: [SEP] Detected Objects:
Ground Truth Exp : nobody likes getting one hour of their life sucked away.
Predicted Exp : the author hates spring.

Original Post : 'can anyone of you imagine a better way to start the new week than having a salivary gland biopsy on monday morning ? me neither ... emoji_300' [SEP] Sarcasm Target: having a salivary gland biopsy [SEP] Image Description: [SEP] Detected Objects:
Ground Truth Exp : having a salivary gland biopsy on monday morning is not a good way to start the new week.
Predicted Exp : having a salivary gland biopsy on monday morning isn't a good way to start the new week.

Original Post : phew ! it 's going to be scorching hot this w-end ! the high on saturday is - 1 ! windchill will prob be - 30 . emoji_619 emoji_300 [SEP] Sarcasm Target: freezing weather [SEP] Image Description: [SEP] Detected Objects:
Ground Truth Exp : the author is worried that the weekend is going to be freezing with a high of -1 and windchill probably -30.
Predicted Exp : it's going to be freezing weather this w-end, the high on saturday is - 1, windchill will prob be - 30.

References

Task I

- <https://medium.com/data-science/build-your-own-transformer-from-scratch-using-pytorch-84c850470dcb>
- <https://www.kaggle.com/code/arunmohan003/transformer-from-scratch-using-pytorch>
- <https://medium.com/@sayedebad.777/building-a-transformer-from-scratch-a-step-by-step-guide-a3df0aeb7c9a>
- <https://www.mislavjuric.com/transformer-from-scratch-in-pytorch/>

Task II

Task III

- [Mutli Head Attention](#)
- Goel, P., Chauhan, D. S., & Akhtar, M. S. (2025, February 11).

Target-Augmented Shared Fusion-based Multimodal Sarcasm Explanation Generation. arXiv.org. <https://arxiv.org/abs/2502.07391>

Thank you!